

PYQ Conceptual-Linkage Analysis

Phase 2 — Final Report

Targets 2015–2026 (N = 1,200) · Priors 1995–2025

Strict 20-year lookback · No future-to-past leakage

7%

Conservative

L4+ directly answerable
~7 questions / paper

20%

Moderate

L3+ mechanism-level help
~20 questions / paper

46%

Broad

L2+ framework familiarity
~46 questions / paper

What this report is

- ▶ Expert conceptual (not lexical) analysis of UPSC Prelims PYQs, 1995–2026.
- ▶ Level 0–5 relationship annotation + advantage types for target → prior links.
- ▶ Design-weighted census: 349 expert overrides on 353 labeled targets.
- ▶ Engine: run13 frozen. Standard: honest numbers, traceable to real questions.
- ▶ Full §1–§24 in the brief's deliverable order, plus Appendix B timelines.

Appendix A (1,200 question-level links) ships as CSV alongside this PDF.

Machine-readable tables: data/pyq_phase2_calibrated_*.json

UPSC Prelims PYQ Conceptual-Linkage Analysis (Phase 2) — Final Report

Targets: GS Paper 1, 2015–2026 (N = 1200). **Priors:** 1995–2025 with strict 20-year lookback (max observed lag = 20; zero links exceed it) and strict no-future-leakage (every prior year < target year, machine-asserted). **Engine:** run13 FROZEN ([scripts/pyq_phase2.py](#)). **Expert review:** 349 overrides, 353 design-labeled targets (173 census + 180 random), 26 targeted extras, 30 rareopt2 pair-checks. Standard: honest, traceable-to-real-questions numbers.

§-mapping note: the brief asked for §24 order + Appendices A/B. No §24 template file exists in the repo, so the 24 sections below follow the brief's deliverable list in logical order. §14/§15 tables are excerpts; the full machine-readable tables ship alongside this report (§24).

§1 Objective & scope

Measure, **conceptually rather than lexically**, how much of each 2015–2026 UPSC Prelims GS paper is answerable/helped by prior-year questions (PYQs) within a 20-year lookback. Outputs: L0–L5 relationship annotation + advantage types for target→prior links; knowledge graph; transformation taxonomy with frequencies; Top 100 high-leverage PYQs; ranked recurring-concept map; year-by-year and subject-wise coverage with 7/10/15/20-yr marginal value; signal-vs-noise for old PYQs; conservative/moderate/broad coverage answers; subject PYQ-depth recommendations; a PYQ→Concept study system; Immortal vs Dead concepts; and future-paper pattern inference (not prediction) — all with confidence labels.

§2 Method (engine + expert calibration)

- 1. Engine (run13, frozen).** Every 2015–2026 target gets ranked prior-links with features (focus-f, shared entities/rare-terms/clusters, option-overlap) and an auto tier L0–L5. Population: auto-L5: 15, L4: 115, L3: 564, L2: 43, L1: 270, L0: 193.
- 2. Expert calibration (design-based).** Census of auto-L5 (15/15), auto-L4 (115/115), auto-L2 (43/43) + simple random samples of auto-L3 (80/564), auto-L1 (50/270), auto-L0 (50/193). 353 labeled targets; each labeled target carries a design weight (1.0 for census tiers; 7.05 / 5.4 / 3.86 for the L3/L1/L0 samples). All population figures are design-weighted; year/subject tables use cell (tier×year, tier×subject) weighting. Only 3 of 120 cells needed tier-mean borrowing (flagged in tables).
- 3. Targeted extras (26).** Early-review finds in auto-L3/L1 (upward-biased by construction): **excluded** from all weighted estimates, **included** in verified counts.
- 4. Link selection.** Expert best-link = engine best-link unless the expert found a better ALT (then link-switched) or found a missed prior by corpus search (27 engine-missed links, mostly in the auto-L0 sample). Every override records tier + link + reason ([data/pyq_phase2_overrides.json](#), 349 entries).

§3 Level definitions & boundary rules (expert)

- L5** — (near-)verbatim repeat: prior determines the answer (e.g. 2021_P_11 × 2019_P_90, money-multiplier, identical).

- **L4** — partial-direct overlap: prior answers ≥ 1 statement/option or refutes one directly (e.g. 2026_P_36 boundary-counting excluded; 2025_P_10 capital-receipts $\times$ 2016_P_12 capital-budget).
- **L3** — same mechanism/Act/scheme, different provision, effect, or chemical (e.g. PFAS $\times$ brominated flame-retardants — same persistence+bioaccumulation logic; GEAC-under-EPA $\times$ GEAC-under-ministry).
- **L2** — framework familiarity: prior teaches the framework/topic but not the tested fact (e.g. Bollgard $\times$ Cas9; Sangam $\times$ Mrichchhakatika-scholars).
- **L1** — theme/format familiarity only (e.g. MDR $\times$ import-cover; town-pairs $\times$ region-pairs).
- **L0** — no meaningful link (incl. false friends: red-tides/algae $\times$ ocean-tides; Krishna-river $\times$ Lord-Krishna; summit-preamble $\times$ Constitution-preamble; Bitcoin $\times$ Fanam-coins; Surat-city $\times$ Surat-split).

Boundary rules applied: partial-direct statement overlap = L4; same-Act-different-provision / same-mechanism-different-effect = L3 (explanatory-bridge test decides); framework-without-fact = L2; same-subject theme/format-only = L1 (generous by design); cross-subject word-only = L0. Zero conflicts between early verdicts and strata verdicts on overlapping questions.

§4 Calibrated census (headline numbers)

Design-weighted estimates for $N = 1200$ (95% Wilson CIs, census tiers exact):

Tier	Est. count	Share	95% CI (count)
L5 (verbatim)	13	1.1%	— (verified 15; upper 72 reflects unsampled mass)
L4+ (answerable)	85	7.1%	[57, 167]
L3+ (mechanism+)	235	19.6%	[174, 339]
L2+ (framework+)	549	45.7%	[449, 666]
L1+ (any signal)	981	81.8%	[870, 1061]
L0 (no link)	219	18.2%	—

Full split: $L5 \approx 13$, $L4 \approx 72$, $L3 \approx 150$, $L2 \approx 314$, $L1 \approx 433$, $L0 \approx 219$. **Verified** (seen by expert): 15 L5 + 35 L4 = 50 L4+; the estimate (85) projects sampled rates onto unsampled mass. Honest reading: roughly **1 in 14 questions is directly answerable from a PYQ; 1 in 5 shares mechanism; nearly half share framework; 4 in 5 carry at least theme signal.**

§5 Engine diagnostics (auto vs expert)

Auto tier	n	P(true L1+)	P(true L2+)	P(true L3+)	P(true L4+)	P(true L5)
L5 (census)	15	1.00	1.00	0.93	0.73	0.47
L4 (census)	115	1.00	0.95	0.59	0.18	0.05
L3 (sample)	80	0.89	0.57	0.21	0.07	0.00
L2 (census)	43	0.84	0.40	0.12	0.07	0.00
L1 (sample)	50	0.78	0.18	0.06	0.00	0.00
L0 (sample)	50	0.54	0.18	0.06	0.04	0.00

Takeaways: (a) auto-L4 over-fires badly as "L4" (18%) but is an excellent L2+ screen (95%) — treat auto-L4 as "needs expert triage", not as truth; (b) auto-L3 precision for true L3+ is only 21%, but

57% carry L2+ signal; (c) auto-L0 misses weak links 54% of the time (mostly L1 theme links found only by active corpus search) — engine-L0 $\neq$ true-L0; (d) upgrades (missed L4+) hide in every tier: auto-L3 7.5%, auto-L2 7%, auto-L1 0% (sample), auto-L0 4%.

§6 Year-by-year conceptual coverage

Cell-weighted shares (% of each year's 100 Qs at/above tier) + mean lag of L2+ links:

Year	L4+	L3+	L2+	L1+	mean lag
2015	8.0	28.7	33.8	62.5	6.3
2016	5.0	6.0	46.7	89.1	4.3
2017	5.0	33.1	40.1	83.8	5.8
2018	2.0	12.6	21.1	81.0	9.3
2019	1.0	8.0	68.7	75.0	8.0
2020	19.9	31.5	59.0	89.5	7.1
2021	5.0	18.0	39.8	82.6	6.2
2022	8.8	24.2	51.8	83.7	8.3
2023	9.0	21.0	66.3	75.0	7.7
2024	5.0	19.5	69.5	83.0	8.9
2025	2.0	11.3	29.9	70.3	7.8
2026	0.0	7.9	43.5	85.1	8.8

Notes: 2020 is the repeat-peak year (L4+ $\approx$ 20%: Welfare-State, money-multiplier-in-2021's-prior, fertigation, sugarcane...). Year cells are small — single-year L3+/L4+ figures are volatile (e.g. 2016 L3+ 6% vs 2017 33%); trust the pooled 19.6% and the block trends (§19), not single-year spikes.

§7 Subject-wise conceptual coverage

Cell-weighted (pop = target counts 2015–2026):

Subject	pop	L4+	L3+	L2+	L1+	mean lag	lag80
History	149	12.6	27.6	57.4	86.0	6.1	6
Economy	239	5.6	29.4	56.3	78.7	8.8	16
Science & Tech	153	12.1	16.9	48.1	83.1	6.1	8
Geography	199	5.6	17.9	46.4	77.7	7.4	12
Art & Culture	56	1.8	3.6	45.4	90.6	9.6	13
Environment	100	7.7	18.3	50.3	92.4	6.6	9
Polity	302	5.9	15.9	29.0	79.1	6.8	11

History and S&T have the densest strong-link mass (L4+ $\approx$ 12–13%); Polity is the weakest (L2+ only 29% — Polity re-tests fresh provisions yearly); Art & Culture is nearly all weak familiarity (L3+ 3.6%, L1+ 91%) — theme helps, facts don't repeat. Env-auto0 and S&T-auto1 cells borrowed tier means (flagged, negligible effect).

§8 Window analysis: 7/10/15/20-yr marginal value

Share of ALL 1200 targets whose best-link falls within lag $\leq W$ (weighted %):

Window	L4+	L3+	L2+ (marginal)	L1+
3 yr	2.4	4.7	11.8	30.4
5 yr	4.2	8.0	18.7 (+6.8)	42.7
7 yr	5.1	11.0	26.2 (+7.5)	51.9
10 yr	6.3	15.3	35.4 (+9.3)	63.9
15 yr	6.8	17.5	41.1 (+5.6)	73.0
20 yr	7.1	19.6	45.7 (+4.6)	81.8

The 5→10-yr band is the richest marginal zone (+16.8 L2+ points); 10→20 yrs still adds +10.3. A 10-year PYQ file captures ~3/4 of the 20-year L2+ value (35.4/45.7) and ~4/5 of L3+ value (15.3/19.6) — the standard "10-year PYQ" advice is roughly right for strong links, but the second decade carries real weight for framework familiarity.

§9 Conservative / moderate / broad 20-year coverage answer

- **Conservative (directly answerable, L4+): $\approx 7\%$ — about 7 questions per paper** (CI 5–14%).
- **Moderate (mechanism-level help, L3+): $\approx 20\%$ — about 20 questions per paper** (CI 15–28%).
- **Broad (framework/theme help, L2+/L1+): $\approx 46\%$ / 82%** (CIs 37–56% / 73–88%).

Recommendation: quote the moderate figure (1 in 5) as "PYQs substantially help", the conservative (1 in 14) as "PYQs directly answer", and never quote L1+ (82%) without the "theme-only" qualifier.

§10 Advantage types

Primary advantage per target (constructed mapping from expert tier + option-overlap; documented, not independently expert-verified):

Advantage	Weighted Qs	Share	Meaning
answerability	85	7.1%	prior (near-)determines the answer (L4–L5)
conceptual	300	25.0%	prior teaches the tested mechanism/framework (L3, most L2)
elimination	164	13.7%	prior kills ≥ 1 distractor (L2–L3 with option-overlap)
pattern	433	36.1%	theme/format familiarity only (L1)
none	219	18.2%	(L0)

Among linked L2+ questions: 15.5% answerability / 54.6% conceptual / 29.9% elimination. Multi-label reality: most L3–L4 links also eliminate; the table assigns each link its primary edge.

§11 Transformation taxonomy (with frequencies)

Constructed classifier on link features (weighted over 1200; expert labels where present):

Class	W	What it is	Example
entity-bridge	278	same entity/concept, new question	GEAC-ministry → GEAC-statute
mechanism-transfer	114	same mechanism, different domain	BFRs → PFAS
angle-shift	45	same theme, different facet	OMO-RBI-tool → OMO-in-MP
same-fact-new-frame	20	same fact, new frame	Swadeshi <-> Partition
paraphrase-repeat	19	reworded repeat	Money-Bill-NOT variants
verbatim-repeat	4	$f \geq 0.7$ near-identical	Buddha-kingdoms 2014→2015
option-echo	0	distractor overlap only, no entity/cluster	(never occurs pure in L2+)
surface-echo	23	thin lexical resemblance	FOBS × asteroids
weak-association	363	L1 with features	town-pairs × region-pairs
expert-verified-weak/strong	93/12	expert-found, no engine features	Ashoka × Kanganahalli
unclassified-weak	11	residual	—
spurious	219	L0	red-tides × ocean-tides

(Entity-bridge dominates strong links; mechanism-transfer is the prized L3 class; pure option-echo without entity/cluster support does not occur — distractor overlap always rides a real concept.)

§12 Signal vs noise: old PYQs

L2+ share of best-links by prior cohort (weighted; w = cohort link-mass):

Prior cohort	w	L4+	L3+	L2+
1995–1999	17	6.0	23.8	35.7
2000–2004	77	7.8	32.5	58.6
2005–2009	70	4.3	23.3	60.9
2010–2014	348	8.1	28.8	68.1
2015–2019	326	13.8	23.1	46.8
2020–2025	144	1.4	9.7	45.4

Old PYQs are NOT noise: 2000–2009 priors deliver L2+ $\approx$ 60% — but the mass is thin (w = 147 vs 348 for 2010–14). The 2010–14 block is the golden corpus (68% L2+, 29% L3+). Strongest verbatims come from 2015–19 priors (13.8% L4+ — adjacent-year repeats). Noise concentrates in: (a) pre-2000 thin mass, (b) cross-subject word-matches (the L0 false-friend catalogue, §3), (c) UNFCCC/GEF-type "theme hubs" that link everywhere at L1 (international-orgs: 14/19 links L1).

§13 Subject PYQ-depth recommendations

Optimal lookback = lag80 of L2+ links (captures 80% of framework-level value):

Subject	Depth	Rationale
Economy	16 yrs	deepest memory (monetary-policy canon from 2000s still pays)
Art & Culture	13 yrs	weak links only — read for theme, not facts
Geography	12 yrs	place/concept links decay slowly
Polity	11 yrs	provisions churn; depth helps procedure/federalism only
Environment	9 yrs	conventions/protocols recur; facts update

Subject	Depth	Rationale
Science & Tech	8 yrs	tech churns; biotech/health mechanisms persist
History	6 yrs (+canon)	80% of value is recent — BUT the deep tail (lags 15–19) is all L3–L5 freedom-movement canon (Surat, 1813, INM-chrono). Do 6 yrs + the canon list, not 20 yrs flat.

§14 Top 100 high-leverage PYQs

Ranked by distinct-target best-link count over the FULL 1200-question graph (732 priors), with expert-quality annotations where labeled. Top 40 (full 100: [data/pyq_phase2_calibrated_top100.json](#)):

Prior	n	ALT	labeled/expert-best	Span	Subject
2010_P_55 (UNFCCC treaty)	7	8	2 / L2	2015–2026	pre2015
2017_P_43 (UPI)	7	6	1 / L3	2018–2026	Economy
2013_P_23 (FRA authority)	7	3	3 / L3	2018–2026	pre2015
2016_P_84 (DigiLocker)	6	4	1 / L1	2017–2023	Polity
2010_P_20 (malaria vaccine)	6	4	0 / —	2019–2026	pre2015
2013_P_61 (kharif season)	6	3	1 / L3	2019–2024	pre2015
2012_P_4 (NGT/Art-21)	5	6	2 / L2	2016–2024	pre2015
2011_P_48	5	5	0 / —	2015–2018	pre2015
2014_P_24 (steel pollutants)	5	5	2 / L3	2016–2021	pre2015
2004_P_78 (foodgrain costs)	5	3	1 / L3	2016–2020	pre2015
2017_P_56 (GIF)	5	3	0 / —	2019–2026	Economy
2013_P_54 (Ajanta)	5	2	1 / L2	2016–2026	pre2015
2011_P_32	5	1	0 / —	2017–2025	pre2015
2014_P_6 (GEF)	4	6	3 / L1	2015–2022	pre2015
2021_P_13 (G-Sec)	4	6	1 / L3	2023–2025	Economy
2004_P_93 (nuclear roles)	4	4	1 / L2	2015–2019	pre2015
2014_P_46 (wetlands)	4	4	2 / L2	2015–2023	pre2015
2017_P_76 (elections)	4	4	2 / L2	2018–2022	Polity
2015_P_1 (PMJDY)	4	2	3 / L3	2016–2024	Economy
2011_P_75 (mangroves)	4	2	2 / L2	2016–2026	pre2015
2015_P_79 (Fortaleza)	4	2	1 / L1	2016–2025	Polity
2010_P_35 (H2 exhaust)	4	1	3 / L5	2015–2025	pre2015
2003_P_42 (Mrichchhakatika)	4	1	2 / L1	2020–2022	pre2015
2021_P_6 (bonds)	4	1	1 / L2	2022–2025	Economy
2003_P_96	4	0	0 / —	2016–2018	pre2015
2016_P_56 (FSDC)	4	0	1 / L1	2017–2026	Economy
2007_P_7 (91st Amdt)	3	8	2 / L4	2018–2023	pre2015

Prior	n	ALT	labeled/expert-best	Span	Subject
2002_P_21 (WIPO)	3	5	1 / L1	2015–2017	pre2015
2002_P_17	3	5	0 / —	2017–2022	pre2015
2015_P_18 (fly ash)	3	5	1 / L2	2018–2020	Environment
2012_P_57 (joint sitting)	3	4	2 / L3	2016–2023	pre2015
2004_P_41	3	4	0 / —	2016–2023	pre2015
2007_P_6	3	4	1 / L1	2019–2023	pre2015
2020_P_65 (G20)	3	4	1 / L1	2022–2024	Economy
2014_P_78 (ITPGRFA)	3	3	0 / —	2015–2016	pre2015
2012_P_84 (coal import)	3	3	1 / L5	2015–2019	pre2015
2013_P_95 (NDC)	3	3	0 / —	2016–2025	pre2015
2013_P_5 (recombinant DNA)	3	3	1 / L2	2016–2022	pre2015
2010_P_45 (Stiglitz Commission)	3	3	0 / —	2017–2022	pre2015
2017_P_91 (Schedule I wildlife)	3	3	1 / L3	2019–2022	Polity

Hubs like UNFCCC/GEF/DigiLocker rank high on count but expert-best is L1–L2 — they are theme magnets, not answer keys. High-leverage quality priors are the L4–L5 ones with smaller n (H2-exhaust, 91st Amendment, coal-import/steel). Full 100: [data/pyq_phase2_calibrated_top100.js on](#).

§15 Ranked recurring-concept map

75 concepts from 323 labeled best-links that carry engine cluster features (27 expert-found links have no engine features and are excluded from this map). **High-confidence** rule (pre-registered): ≥ 4 links, ≥ 3 distinct target-years, ≥ 1 L3+. **Low-confidence / thin** = the rest.

Band	n	How to read
High-confidence	35	Recurs across years at mechanism-or-better; worth a dedicated PYQ sheet
Low-confidence / thin	40	Either theme-only (protected-areas, judiciary) or 1–2 isolated hits (Cold concepts, §18)

Top high-confidence concepts by design weight (full timelines: Appendix B):

Concept	w	links	years	best expert tier	tag
international-orgs	65	19	6	L4	IMMORTAL — but 14/19 are L1 theme-hubs
minerals-energy	46	10	7	L3	IMMORTAL
soils-agrogeo	40	10	6	L4	IMMORTAL
biotech-health	32	8	6	L4	IMMORTAL
monetary-policy	31	7	5	L5	IMMORTAL
fundamental-rights	30	12	7	L5	IMMORTAL
statutory-bodies	29	5	5	L3	IMMORTAL
chemistry-funda	27	9	6	L5	IMMORTAL
rivers-drainage	27	9	7	L3	IMMORTAL
buddhism-jainism	25	7	6	L5	IMMORTAL

Concept	w	links	years	best expert tier	tag
capital-markets	24	6	6	L3	IMMORTAL
parliament-structure	23	11	8	L4	IMMORTAL
biology-funda	23	5	5	L4	IMMORTAL
ecology-basics	22	5	4	L4	IMMORTAL
env-laws	18	6	6	L3	high-conf, not immortal (span)

Caveat: [international-orgs](#) is the heaviest node and the noisiest — treat as a theme list (UNFCCC, GEF, WIPO, G20), not as a mechanism to memorise. Monetary-policy, fundamental-rights, parliament-structure, and env-laws are the high-confidence *and* high-quality set.

§16 Knowledge graph

Directed graph of target → best-prior, plus 75 concept nodes.

Layer	Count	Notes
Target nodes (2015–2026)	1,200	one per GS-I question
Prior nodes with ≥1 best-link	732	from the full auto graph
Concept nodes	75	shared-cluster labels
Best-link edges	1,200	one per target (None if L0)
Expert-labeled edges	353	census + samples; design-weighted
Verified L5 edges	15	listed in §4
Verified L4 edges	35	listed in §4

Machine-readable edges: [data/pyq_phase2_calibrated_edges.json](#). Concept payloads: [data/pyq_phase2_calibrated_concepts.json](#). The graph is **not** a citation network of equal-weight edges — an L5 edge and an L1 edge are different kinds of thing; always filter on expert tier.

§17 PYQ → Concept study system

Do **not** read PYQs year-by-year from 1995. Use this loop:

- Pick an Immortal concept** from §18 (start with monetary-policy, fundamental-rights, parliament-structure/procedure, env-laws, soils-agrogeo, biotech-health).
- Read the concept's example pairs** in Appendix B (target × prior, with expert tier). For every L3+ pair, write one sentence: *what mechanism transferred?*
- Add the high-leverage priors** in §14 that sit inside that concept (e.g. monetary-policy ← 2010_P_35 H2-exhaust, 2013_P_72, 2011_P_87).
- Stop at lag80 for the subject** (§13). Extra years only for the History freedom-movement canon.
- Drill L4-L5 pairs as flashcards** (answerability). Drill L3 pairs as "same Act / new provision" or "same mechanism / new domain". Ignore L1 pairs except as a theme reminder.

Time budget that matches the numbers: ~80% of PYQ hours on the 35 high-confidence concepts; ~15% on subject-depth tails (Economy 16-yr, Geography 12-yr); ~5% on the History canon list. Do not farm international-orgs / DigiLocker / GEF hubs for facts.

§18 Immortal vs Dead concepts

Rule (pre-registered, Appendix B): **Immortal** = span ≥ 8 years with ≥ 3 L2+ links. **Cold / Dead** = tagged when the concept has 1–2 isolated hits and does not recur.

Immortal (25): international-orgs, minerals-energy, soils-agrogeo, biotech-health, monetary-policy, fundamental-rights, statutory-bodies, chemistry-funda, rivers-drainage, buddhism-jainism, capital-markets, biology-funda, parliament-structure, ecology-basics, indian-physiography, union-executive, banking, parliament-procedure, congress-early, dpsp-duties, climate-change, world-geo, state-executive, federalism, scheme-finclusion.

These are the concepts UPSC has kept returning to across a decade. They are not "will appear next year" predictions — they are the ones where a 10–16 year PYQ file still pays.

Cold / Dead (20): post-mauryan, harappan, mughals, gupta-harsha, india-foreign-policy, scheme-agri, awards-sports-gk, vedic-mahajanapada, national-income, transport-ports, elections, emergency, nuclear, employment-poverty, industry-services, reports-indices, scheme-food-nutrition, personalities-books, govt-programmes-misc, oceanography.

Cold $\neq$ unimportant for GS. It means PYQ-to-PYQ transfer inside 2015–2026 is thin. A Harappan fact can still be asked; a prior Harappan PYQ is unlikely to answer it. Do not skip the static syllabus; skip the false hope that old PYQs will decode the next question.

§19 Future-paper pattern inference (not prediction)

Cell-weighted block means from §6 (each year = 100 Qs; not a forecast):

Block	L4+	L3+	L2+	L1+	Reading
2015–2018	5.0	20.1	35.4	79.1	mechanism mass present, few verbatims
2019–2022	8.7	20.4	54.8	82.7	2020 spike pulls L4+ up; L2+ jumps
2023–2026	4.0	14.9	52.3	78.4	fewer direct repeats; framework help holds

What this licenses, and what it does not:

- **Expect ~5–9 L4+ and ~15–25 L3+ per paper**, not a rising "UPSC repeats more now" story. 2020 is an outlier (L4+ $\approx$ 20%), 2026 sits at the floor (L4+ = 0 in the cell-weighted estimate).
- **L2+ around half the paper is the stable band** (2019–2026). Theme familiarity is the median help, not the exception.
- **Polity will keep looking "unlinked"** (L2+ 29%): new provisions, not recycled ones. History/S&T will keep looking "linked" (L4+ $\sim$ 12%).
- **Adjacent-year verbatim repeats happen and are rare** (L5 $\approx$ 1%). Do not study last year's paper as an answer key; do study it as the newest prior.
- **Mechanism-transfer (the L3 class) is the durable skill** — PFAS $\leftarrow$ BFRs, GEAC-statute $\leftarrow$ GEAC-ministry, capital-receipts $\leftarrow$ capital-budget. That is the pattern to practise, not a topic list to guess from.

No year, subject, or concept is a prediction of 2027 contents.

§20 Confidence labels

Claim	Confidence	Why
Auto-tier precision table (§5)	High	Census of auto-L5/L4/L2; sample n = 80/50/50 for L3/L1/L0
Pooled L4+ ≈ 7%, L3+ ≈ 20%, L2+ ≈ 46%	High	Design-weighted; Wilson CIs reported
Single-year L3+/L4+ (e.g. 2016 vs 2017)	Low	100-Q cells; one upgrade swings the year
Subject L4+/L3+ ordering (Hist/S&T dense; Polity thin)	Moderate-high	Cell-weighted; 3/120 cells borrowed tier means
Window 5→10 yr as richest marginal band	High	Same weights as pooled census
Advantage-type shares (§10)	Moderate	Constructed from tier + option-overlap; not independently labeled
Taxonomy frequencies (§11)	Moderate	Constructed classifier; expert labels used where present
Top-100 *count* ranking	High as a hub list; Low as a quality list	Count ≠ expert tier (UNFCCC is a hub, not an answer key)
Immortal / Cold tags	Moderate	Rule-based on labeled cluster-links only (323 of 353)
Block trend 2023–26 "fewer repeats"	Moderate	Four years, 2026 at floor; do not overfit
"Next paper will look like X"	None	Explicitly not claimed

§21 Limitations & caveats

- **Engine is a screen, not a judge.** Auto-L4 precision for true L4 is 18%; auto-L0 is not true-L0 (54% hide at least L1). Every population number in this report is expert-calibrated, not engine-raw.
- **27 engine-missed links** were found by active corpus search, mostly in the auto-L0 sample. The L0→L1 leak is a lower bound on missed weak links.
- **Targeted extras (26)** are excluded from weighted estimates (upward-biased by construction) and included in verified L4+/L5 counts. Mixing the two is the usual way to over-claim.
- **Subject tags** are the corpus tags; a few questions sit on borders (agri-env, scheme-finclusion). 0 ther is 2 questions and is ignored.
- **2026** is one paper. Its L4+ = 0 is a point estimate, not a regime change.
- **Pre-2015 priors** are in the graph as nodes but were not themselves expert-tiered as targets. Quality annotations on old hubs exist only when a 2015–2026 labeled target pointed at them.
- **No Mains, no CSAT, no BPSC.** This is GS Paper I Prelims only.

§22 How an aspirant should use this

- Quote **1 in 5** (L3+) when someone asks "do PYQs help?", **1 in 14** (L4+) when they ask "will this question repeat?", and never quote 82% without saying "theme only".
- Build a **10-year PYQ file as the default**, then add years 11–16 only for Economy and the History freedom-movement canon (Surat, 1813 Charter, INM chronology).
- Study **concepts, not years**: one Immortal concept per sitting, with its L3+ pairs from Appendix B.

- Treat auto-linked "repeats" in coaching handouts as **triage queues**. Most auto-L4s are L2. Verify before memorising.
- Keep a **false-friend list** (red-tides ≠ ocean-tides; Krishna-river ≠ Lord-Krishna; Bitcoin ≠ Fanam; Surat-city ≠ Surat-split). Those are the engine's favourite L0 traps and a good model of how lexical PYQ tools mislead.

§23 Reproducibility

Item	Where
Frozen engine	scripts/pyq_phase2.py (run13)
Target corpus 2015–2026	data/pyq_phase2/GS_YYYY.json (12 files)
349 expert overrides	data/pyq_phase2_overrides.json
Design-labeled links	data/pyq_phase2_calibrated_links.json (353)
Full 1,200-edge graph	data/pyq_phase2_calibrated_edges.json
Top 100	data/pyq_phase2_calibrated_top100.json
75 concepts	data/pyq_phase2_calibrated_concepts.json
Subject estimates	data/pyq_phase2_calibrated_subjects.json
Analysis scripts	scripts/pyq_phase2_{analyze,cells,coverage,taxonomy,top100,appB}.py

Lookback assertion: max observed lag = 20; every prior year < target year (machine-asserted).
 Uncalibrated leftovers `data/pyq_phase2_{summary,top300,concepts}.json` are **not** sources of truth.

§24 Deliverable index

#	Deliverable	File
1–23	This report	docs/PYQ_PHASE2_REPORT.md / .pdf
A	All 1,200 question-level links	docs/PYQ_PHASE2_APPENDIX_A.csv
B	75 concept timelines	docs/PYQ_PHASE2_APPENDIX_B.md
—	HTML viewer (UPSC section)	preview.html (nav: UPSC → Phase 2 Report)

Standard held: honest, traceable-to-real-questions numbers over impressive ones.

Appendix B — Concept timelines

75 concept timelines from the calibrated knowledge graph. Full question-level links (Appendix A) are in `docs/PYQ_PHASE2_APPENDIX_A.csv`.

Appendix B — Concept timelines (shared-cluster links, labeled targets)

Covers 75 concepts from 323 labeled best-links with engine features (27 expert-found links lack engine features and are excluded here). Tiers are EXPERT tiers. High-confidence = ≥ 4 links, ≥ 3 distinct target-years, ≥ 1 L3+. Immortal = span $\geq 8y$ with ≥ 3 L2+ links.

B1. High-confidence recurring concepts

international-orgs [IMMORTAL]

- weight=65 links=19 priors=17 years=2015–2025 (6 distinct) tiers={1: 14, 2: 3, 3: 1, 4: 1}
- examples: 2015_P_57(L4)x2008_P_61; 2015_P_79(L1)x2009_P_47; 2016_P_51(L1)x2015_P_43

minerals-energy [IMMORTAL]

- weight=46 links=10 priors=10 years=2015–2026 (7 distinct) tiers={1: 1, 2: 6, 3: 3}
- examples: 2015_P_8(L2)x2003_P_36; 2016_P_74(L2)x2014_P_38; 2018_P_27(L2)x2008_P_47

soils-agrogeo [IMMORTAL]

- weight=40 links=10 priors=9 years=2016–2026 (6 distinct) tiers={1: 3, 2: 4, 3: 2, 4: 1}
- examples: 2016_P_40(L1)x2012_P_4; 2019_P_21(L3)x2012_P_81; 2019_P_83(L3)x2013_P_61

biotech-health [IMMORTAL]

- weight=32 links=8 priors=7 years=2015–2026 (6 distinct) tiers={1: 3, 2: 3, 3: 1, 4: 1}
- examples: 2015_P_63(L3)x2003_P_19; 2019_P_100(L2)x2013_P_43; 2020_P_41(L4)x2015_P_23

monetary-policy [IMMORTAL]

- weight=31 links=7 priors=7 years=2015–2024 (5 distinct) tiers={1: 1, 3: 4, 4: 1, 5: 1}
- examples: 2015_P_22(L3)x2011_P_87; 2015_P_86(L4)x2013_P_72; 2020_P_70(L3)x2012_P_75

fundamental-rights [IMMORTAL]

- weight=30 links=12 priors=9 years=2015–2025 (7 distinct) tiers={1: 2, 2: 4, 3: 2, 4: 1, 5: 3}
- examples: 2015_P_83(L2)x2013_P_25; 2015_P_89(L4)x2013_P_25; 2017_P_20(L1)x2016_P_20

statutory-bodies [IMMORTAL]

- weight=29 links=5 priors=5 years=2015–2026 (5 distinct) tiers={1: 1, 3: 4}
- examples: 2015_P_22(L3)x2011_P_87; 2020_P_70(L3)x2012_P_75; 2023_P_24(L3)x2010_P_80

chemistry-funda [IMMORTAL]

- weight=27 links=9 priors=7 years=2015–2026 (6 distinct) tiers={1: 2, 2: 6, 5: 1}
- examples: 2015_P_8(L2)x2003_P_36; 2016_P_24(L2)x2013_P_31; 2023_P_60(L2)x2010_P_35

rivers-drainage [IMMORTAL]

- weight=27 links=9 priors=9 years=2016–2024 (7 distinct) tiers={1: 2, 2: 6, 3: 1}

- examples: 2016_P_85(L2)x2003_P_78; 2017_P_2(L3)x2011_P_91; 2017_P_30(L2)x2011_P_96

buddhism-jainism [IMMORTAL]

- weight=25 links=7 priors=6 years=2015-2026 (6 distinct) tiers={2: 5, 4: 1, 5: 1}
- examples: 2015_P_39(L5)x2014_P_25; 2017_P_69(L2)x2010_P_54; 2020_P_22(L4)x2019_P_9

capital-markets [IMMORTAL]

- weight=24 links=6 priors=5 years=2018-2026 (6 distinct) tiers={1: 1, 2: 3, 3: 2}
- examples: 2018_P_56(L3)x2001_P_48; 2020_P_54(L2)x2001_P_48; 2022_P_65(L2)x2021_P_6

biology-funda [IMMORTAL]

- weight=23 links=5 priors=5 years=2016-2026 (5 distinct) tiers={1: 1, 2: 3, 4: 1}
- examples: 2016_P_31(L2)x2013_P_5; 2020_P_41(L4)x2015_P_23; 2021_P_73(L1)x2019_P_99

parliament-structure [IMMORTAL]

- weight=23 links=11 priors=9 years=2015-2024 (8 distinct) tiers={1: 2, 2: 2, 3: 4, 4: 3}
- examples: 2015_P_27(L4)x2013_P_1; 2016_P_95(L3)x2006_P_13; 2018_P_1(L1)x2017_P_76

ecology-basics [IMMORTAL]

- weight=22 links=5 priors=4 years=2018-2026 (4 distinct) tiers={1: 2, 2: 1, 3: 1, 4: 1}
- examples: 2018_P_82(L1)x2011_P_47; 2019_P_30(L2)x2014_P_98; 2021_P_22(L4)x2014_P_80

indian-physiography [IMMORTAL]

- weight=18 links=6 priors=6 years=2016-2026 (5 distinct) tiers={1: 1, 2: 3, 3: 1, 4: 1}
- examples: 2016_P_22(L2)x1999_P_48; 2017_P_4(L3)x2008_P_22; 2017_P_58(L2)x2010_P_19

union-executive [IMMORTAL]

- weight=18 links=6 priors=5 years=2017-2025 (4 distinct) tiers={1: 2, 2: 1, 3: 3}
- examples: 2017_P_50(L3)x2013_P_65; 2017_P_96(L3)x2014_P_60; 2020_P_20(L1)x2007_P_6

env-laws

- weight=18 links=6 priors=4 years=2017-2022 (6 distinct) tiers={3: 6}
- examples: 2017_P_91(L3)x2015_P_19; 2018_P_98(L3)x2013_P_23; 2019_P_55(L3)x2013_P_23

pollution-waste

- weight=18 links=6 priors=5 years=2019-2024 (4 distinct) tiers={2: 3, 3: 2, 4: 1}
- examples: 2019_P_35(L3)x2014_P_24; 2019_P_59(L2)x2013_P_26; 2021_P_17(L3)x2014_P_24

banking [IMMORTAL]

- weight=16 links=4 priors=3 years=2015-2026 (4 distinct) tiers={2: 2, 3: 1, 4: 1}
- examples: 2015_P_100(L2)x2010_P_24; 2022_P_64(L4)x2019_P_73; 2023_P_71(L2)x2013_P_73

agri-env

- weight=16 links=4 priors=4 years=2015-2021 (4 distinct) tiers={1: 1, 2: 1, 3: 1, 4: 1}
- examples: 2015_P_63(L3)x2003_P_19; 2018_P_59(L1)x2014_P_1; 2020_P_91(L4)x2011_P_57

climatology

- weight=16 links=4 priors=3 years=2019-2024 (4 distinct) tiers={1: 1, 2: 1, 3: 1, 5: 1}
- examples: 2019_P_25(L2)x2011_P_94; 2021_P_58(L3)x2015_P_66; 2023_P_62(L1)x2004_P_72

literature-languages

- weight=14 links=4 priors=3 years=2015-2026 (4 distinct) tiers={1: 3, 3: 1}
- examples: 2015_P_75(L3)x2014_P_8; 2020_P_36(L1)x2003_P_42; 2022_P_93(L1)x2003_P_42

key-legislations

- weight=14 links=4 priors=2 years=2017-2021 (4 distinct) tiers={3: 4}
- examples: 2017_P_24(L3)x2007_P_11; 2018_P_98(L3)x2013_P_23; 2019_P_55(L3)x2013_P_23

constitutional-reforms

- weight=14 links=8 priors=8 years=2016-2023 (7 distinct) tiers={1: 1, 2: 2, 3: 4, 5: 1}
- examples: 2016_P_64(L5)x2004_P_29; 2017_P_26(L2)x2015_P_28; 2017_P_62(L3)x2012_P_27

parliament-procedure [IMMORTAL]

- weight=13 links=7 priors=6 years=2015-2024 (6 distinct) tiers={2: 1, 3: 4, 4: 2}
- examples: 2015_P_27(L4)x2013_P_1; 2015_P_37(L3)x2003_P_88; 2018_P_31(L3)x2000_P_53

congress-early [IMMORTAL]

- weight=11 links=5 priors=5 years=2015-2023 (3 distinct) tiers={3: 1, 4: 3, 5: 1}
- examples: 2015_P_48(L4)x2000_P_17; 2015_P_68(L3)x1998_P_25; 2016_P_49(L5)x2010_P_82

dpsp-duties [IMMORTAL]

- weight=11 links=11 priors=5 years=2015-2025 (5 distinct) tiers={2: 5, 3: 3, 4: 2, 5: 1}
- examples: 2015_P_7(L4)x2002_P_43; 2015_P_83(L2)x2013_P_25; 2015_P_89(L4)x2013_P_25

external-sector

- weight=10 links=4 priors=3 years=2015-2022 (4 distinct) tiers={2: 1, 3: 3}
- examples: 2015_P_34(L3)x1998_P_83; 2019_P_86(L2)x2002_P_69; 2020_P_51(L3)x2001_P_83

climate-change [IMMORTAL]

- weight=9 links=5 priors=5 years=2015-2026 (4 distinct) tiers={1: 1, 2: 3, 4: 1}
- examples: 2015_P_74(L1)x2014_P_6; 2016_P_30(L4)x2008_P_95; 2016_P_55(L2)x2011_P_14

world-geo [IMMORTAL]

- weight=8 links=8 priors=7 years=2015-2025 (6 distinct) tiers={1: 1, 2: 3, 3: 3, 4: 1}
- examples: 2015_P_64(L3)x2004_P_20; 2017_P_30(L2)x2011_P_96; 2017_P_54(L4)x2015_P_77

preamble-citizenship

- weight=6 links=6 priors=3 years=2015-2020 (3 distinct) tiers={2: 3, 3: 1, 4: 1, 5: 1}
- examples: 2015_P_83(L2)x2013_P_25; 2015_P_89(L4)x2013_P_25; 2017_P_48(L2)x2015_P_89

state-executive [IMMORTAL]

- weight=5 links=5 priors=4 years=2015-2025 (4 distinct) tiers={1: 1, 2: 1, 3: 2, 4: 1}
- examples: 2015_P_82(L1)x2001_P_49; 2017_P_96(L3)x2014_P_60; 2019_P_53(L2)x2007_P_7

federalism [IMMORTAL]

- weight=5 links=5 priors=4 years=2016-2025 (4 distinct) tiers={2: 3, 3: 1, 5: 1}
- examples: 2016_P_95(L3)x2006_P_13; 2020_P_12(L5)x2015_P_89; 2020_P_5(L2)x2015_P_89

british-early

- weight=5 links=5 priors=4 years=2018-2023 (5 distinct) tiers={2: 1, 3: 3, 4: 1}

- examples: 2018_P_68(L3)x2008_P_55; 2019_P_4(L4)x2002_P_59; 2020_P_34(L3)x2018_P_89

biodiversity-hotspots

- weight=4 links=4 priors=4 years=2015-2025 (3 distinct) tiers={1: 2, 2: 1, 3: 1}
- examples: 2015_P_43(L1)x2005_P_45; 2017_P_4(L3)x2008_P_22; 2017_P_58(L2)x2010_P_19

B2. Lower-confidence / thin concepts

early-medieval

- weight=21 links=3 priors=3 years=2023-2026 (3 distinct) tiers={2: 3}
- examples: 2023_P_45(L2)x2017_P_51; 2025_P_14(L2)x2020_P_24; 2026_P_14(L2)x2019_P_11

protected-areas

- weight=20 links=8 priors=7 years=2015-2020 (5 distinct) tiers={1: 3, 2: 5}
- examples: 2015_P_52(L2)x2013_P_11; 2015_P_9(L1)x2014_P_46; 2017_P_95(L1)x2014_P_9

judiciary

- weight=17 links=5 priors=4 years=2016-2021 (4 distinct) tiers={1: 1, 2: 4}
- examples: 2016_P_40(L1)x2012_P_4; 2017_P_92(L2)x1997_P_79; 2019_P_52(L2)x2012_P_4

temple-architecture

- weight=15 links=3 priors=3 years=2015-2026 (3 distinct) tiers={1: 1, 2: 2}
- examples: 2015_P_54(L1)x2001_P_5; 2021_P_49(L2)x2013_P_54; 2026_P_14(L2)x2019_P_11

species-focus

- weight=10 links=4 priors=3 years=2016-2020 (3 distinct) tiers={1: 1, 2: 3}
- examples: 2016_P_22(L2)x1999_P_48; 2017_P_95(L1)x2014_P_9; 2020_P_100(L2)x2012_P_87

inflation

- weight=9 links=3 priors=2 years=2015-2021 (3 distinct) tiers={2: 1, 3: 1, 4: 1}
- examples: 2015_P_29(L3)x2010_P_85; 2020_P_67(L4)x2010_P_85; 2021_P_12(L2)x2011_P_12

forests-vegetation

- weight=9 links=3 priors=3 years=2015-2019 (3 distinct) tiers={1: 1, 2: 2}
- examples: 2015_P_78(L2)x2010_P_68; 2016_P_33(L1)x2012_P_43; 2019_P_33(L2)x2005_P_36

wetlands-mangroves-coral

- weight=9 links=3 priors=2 years=2016-2023 (3 distinct) tiers={1: 2, 2: 1}
- examples: 2016_P_44(L2)x2011_P_75; 2021_P_16(L1)x2019_P_40; 2023_P_5(L1)x2011_P_75

scheme-finclusion [IMMORTAL]

- weight=9 links=3 priors=2 years=2016-2024 (3 distinct) tiers={2: 1, 3: 2}
- examples: 2016_P_73(L2)x2015_P_1; 2017_P_79(L3)x2016_P_4; 2024_P_100(L3)x2015_P_1

it-digital

- weight=9 links=3 priors=3 years=2018-2026 (3 distinct) tiers={1: 2, 3: 1}
- examples: 2018_P_64(L1)x2016_P_84; 2024_P_26(L1)x2022_P_31; 2026_P_55(L3)x2017_P_43

space

- weight=9 links=3 priors=2 years=2023-2025 (2 distinct) tiers={1: 1, 2: 1, 3: 1}
- examples: 2023_P_57(L3)x2018_P_61; 2025_P_83(L1)x2023_P_56; 2025_P_94(L2)x2018_P_61

post-mauryan [COLD]

- weight=8 links=2 priors=1 years=2020-2022 (2 distinct) tiers={1: 2}
- examples: 2020_P_36(L1)x2003_P_42; 2022_P_93(L1)x2003_P_42

physics-funda

- weight=8 links=2 priors=2 years=2020-2022 (2 distinct) tiers={1: 1, 4: 1}
- examples: 2020_P_43(L4)x2017_P_87; 2022_P_34(L1)x2011_P_85

geomorphology

- weight=8 links=2 priors=2 years=2021-2025 (2 distinct) tiers={2: 1, 5: 1}
- examples: 2021_P_62(L5)x2013_P_10; 2025_P_25(L2)x2014_P_13

harappan [COLD]

- weight=7 links=1 priors=1 years=2017-2017 (1 distinct) tiers={3: 1}
- examples: 2017_P_2(L3)x2011_P_91

mughals [COLD]

- weight=7 links=1 priors=1 years=2019-2019 (1 distinct) tiers={2: 1}
- examples: 2019_P_17(L2)x2006_P_75

gupta-harsha [COLD]

- weight=7 links=1 priors=1 years=2020-2020 (1 distinct) tiers={1: 1}
- examples: 2020_P_36(L1)x2003_P_42

india-foreign-policy [COLD]

- weight=7 links=1 priors=1 years=2021-2021 (1 distinct) tiers={3: 1}
- examples: 2021_P_58(L3)x2015_P_66

scheme-agri [COLD]

- weight=7 links=1 priors=1 years=2026-2026 (1 distinct) tiers={1: 1}
- examples: 2026_P_38(L1)x2010_P_48

awards-sports-gk [COLD]

- weight=5 links=1 priors=1 years=2023-2023 (1 distinct) tiers={2: 1}
- examples: 2023_P_95(L2)x2021_P_96

env-conventions

- weight=3 links=3 priors=3 years=2015-2021 (3 distinct) tiers={1: 2, 2: 1}
- examples: 2015_P_43(L1)x2005_P_45; 2018_P_88(L2)x2010_P_55; 2021_P_16(L1)x2019_P_40

fiscal-policy

- weight=3 links=3 priors=3 years=2015-2021 (3 distinct) tiers={4: 1, 5: 2}
- examples: 2015_P_60(L4)x2010_P_76; 2016_P_3(L5)x2015_P_98; 2021_P_10(L5)x2013_P_81

gandhian-movements

- weight=3 links=3 priors=3 years=2017-2026 (3 distinct) tiers={1: 1, 3: 1, 4: 1}

- examples: 2017_P_93(L4)x1998_P_72; 2021_P_43(L1)x2008_P_58; 2026_P_19(L3)x2018_P_69

states-reorganisation

- weight=2 links=2 priors=2 years=2015-2026 (2 distinct) tiers={2: 1, 3: 1}
- examples: 2015_P_59(L3)x2012_P_40; 2026_P_82(L2)x2025_P_56

official-language

- weight=2 links=2 priors=2 years=2015-2024 (2 distinct) tiers={3: 2}
- examples: 2015_P_75(L3)x2014_P_8; 2024_P_72(L3)x2008_P_98

agri-economics

- weight=2 links=2 priors=2 years=2019-2023 (2 distinct) tiers={2: 1, 3: 1}
- examples: 2019_P_79(L3)x2004_P_78; 2023_P_27(L2)x2020_P_69

earth-solar-system

- weight=2 links=2 priors=1 years=2023-2024 (2 distinct) tiers={1: 1, 5: 1}
- examples: 2023_P_62(L1)x2004_P_72; 2024_P_2(L5)x2004_P_72

vedic-mahajanapada [COLD]

- weight=1 links=1 priors=1 years=2015-2015 (1 distinct) tiers={5: 1}
- examples: 2015_P_39(L5)x2014_P_25

national-income [COLD]

- weight=1 links=1 priors=1 years=2015-2015 (1 distinct) tiers={3: 1}
- examples: 2015_P_81(L3)x2002_P_36

transport-ports [COLD]

- weight=1 links=1 priors=1 years=2017-2017 (1 distinct) tiers={1: 1}
- examples: 2017_P_84(L1)x2009_P_54

elections [COLD]

- weight=1 links=1 priors=1 years=2017-2017 (1 distinct) tiers={3: 1}
- examples: 2017_P_90(L3)x2001_P_43

emergency [COLD]

- weight=1 links=1 priors=1 years=2017-2017 (1 distinct) tiers={3: 1}
- examples: 2017_P_96(L3)x2014_P_60

nuclear [COLD]

- weight=1 links=1 priors=1 years=2018-2018 (1 distinct) tiers={2: 1}
- examples: 2018_P_33(L2)x2004_P_93

employment-poverty [COLD]

- weight=1 links=1 priors=1 years=2019-2019 (1 distinct) tiers={3: 1}
- examples: 2019_P_24(L3)x1999_P_10

industry-services [COLD]

- weight=1 links=1 priors=1 years=2019-2019 (1 distinct) tiers={2: 1}
- examples: 2019_P_77(L2)x2016_P_37

reports-indices [COLD]

- weight=1 links=1 priors=1 years=2019-2019 (1 distinct) tiers={2: 1}
- examples: 2019_P_77(L2)x2016_P_37

scheme-food-nutrition [COLD]

- weight=1 links=1 priors=1 years=2019-2019 (1 distinct) tiers={3: 1}
- examples: 2019_P_79(L3)x2004_P_78

personalities-books [COLD]

- weight=1 links=1 priors=1 years=2024-2024 (1 distinct) tiers={1: 1}
- examples: 2024_P_86(L1)x2009_P_61

govt-programmes-misc [COLD]

- weight=1 links=1 priors=1 years=2026-2026 (1 distinct) tiers={3: 1}
- examples: 2026_P_55(L3)x2017_P_43

oceanography [COLD]

- weight=1 links=1 priors=1 years=2026-2026 (1 distinct) tiers={1: 1}
- examples: 2026_P_75(L1)x2006_P_26